

How Learning Differs from Optimization

Sargur N. Srihari
srihari@cedar.buffalo.edu

Topics in Optimization

- Optimization for Training Deep Models: Overview
- How learning differs from optimization
 - Risk, empirical risk and surrogate loss
 - Batch, minibatch, data shuffling
- Challenges in neural network optimization
- Basic Algorithms
- Parameter initialization strategies
- Algorithms with adaptive learning rates
- Approximate second-order methods
- Optimization strategies and meta-algorithms ²

Topics in Learning vs Optimization

- Learning vs Pure Optimization
- Empirical Risk Minimization
- Surrogate Loss Functions and Early Stopping
- Batch and Minibatch Algorithms

Learning vs Pure Optimization

- Optimization algorithms for deep learning differ from traditional optimization in several ways:
 - Machine learning acts indirectly
 - We care about some performance measure P defined wrt the training set which may be intractable
 - We reduce a different cost function $J(\theta)$ in the hope that doing so will reduce P
- Pure optimization: minimizing J is a goal in itself
- Optimizing algorithms for training Deep models:
 - Includes specialization on specific structure of ML objective function

Typical Cost Function

- Cost is average over the training set

$$J(\theta) = E_{(x,y) \sim \hat{p}_{\text{data}}} (L(f(x;\theta), y)) \quad \text{where}$$

- $f(x; \theta)$ is the predicted output when the input is x
 - In supervised learning y is target output
 - L is the per-example loss function
 - \hat{p}_{data} is the empirical distribution
- We consider the unregularized supervised case
 - where arguments of L are $f(x; \theta)$ and y
 - Trivial to extend to cases:
 - Where parameters θ and input x are arguments or
 - Exclude output y as argument
 - For regularization or unsupervised learning

Objective wrt data generation is risk

- Objective function wrt training set is

$$J(\theta) = E_{(x,y) \sim \hat{p}_{\text{data}}} \left(L(f(x;\theta), y) \right)$$

L is the per-example loss function

- We would prefer to minimize the corresponding objective function where expectation is across the data generating distribution p_{data} rather than over finite training set

$$J(\theta) = E_{(x,y) \sim p_{\text{data}}} \left(L(f(x;\theta), y) \right)$$

– The goal of a machine learning algorithm is to reduce this expected generalization error

- This quantity is known as *risk*

Empirical Risk

- True risk is $J(\theta) = E_{(\mathbf{x}, y) \sim p_{\text{data}}} (L(f(\mathbf{x}; \theta), y))$
 - If we knew $p_{\text{data}}(\mathbf{x}, y)$ it would be optimization solved by an optimization algorithm
 - When we do not know $p_{\text{data}}(\mathbf{x}, y)$ but only have a training set of samples, we have a machine learning problem
- Empirical risk, with m training examples, is

$$J(\theta) = E_{(\mathbf{x}, y) \sim \hat{p}_{\text{data}}} (L(f(\mathbf{x}; \theta), y)) = \frac{1}{m} \sum_{i=1}^m L(f(\mathbf{x}^{(i)}; \theta), y^{(i)})$$

Empirical Risk Minimization

- Empirical risk, with m training examples, is

$$J(\theta) = E_{(\mathbf{x}, y) \sim \hat{p}_{\text{data}}} (L(f(\mathbf{x}; \theta), y)) = \frac{1}{m} \sum_{i=1}^m L(f(\mathbf{x}^{(i)}; \theta), y^{(i)})$$

- Which is the average training error
 - Still similar to straightforward optimization
- But *empirical risk minimization* is not very useful:
 1. Prone to overfitting: can simply memorize training set
 2. SGD is commonly used, but many useful loss functions have 0-1 loss, with no useful derivatives (derivative is either 0 or undefined everywhere)
- We must use a slightly different approach
 - Quantity we must optimize is even more different from what we truly want to optimize

Surrogate Loss: Log-likelihood

- Exactly minimizing 0-1 loss is typically intractable (exponential in the input dimension) even for a linear classifier
- In such situations use a *surrogate loss function*
 - Acts as a proxy but has advantages
- *Negative log-likelihood* of the correct class is a surrogate for 0-1 loss
 - It allows model to estimate conditional probability of classes given the input
 - If it does that well then can pick the classes that yield the least classification error in expectation

Surrogate may learn more

- Using log-likelihood surrogate,
 - Test set 0-1 loss continues to decrease for a long time after the training set 0-1 loss has reached zero when training
 - Because one can improve classifier robustness by further pushing the classes apart
 - Results in a more confident and robust classifier
 - Thus extracting more information from the training data than with minimizing 0-1 loss

Learning does not stop at minimum

- Important difference between Optimization in general and Optimization for Training:
 - Training does not halt at a local minimum
 - Early Stopping: Instead Learning algorithm halts on an early stopping criterion
 - Based on a true underlying loss function
 - Such as 0-1 loss measured on a validation set
 - Designed to cause algorithm to stop overfitting
- Often stops when derivatives are still large
 - In pure optimization, algorithm considered to converge when derivatives are very small

Decomposition: Batch Algorithms

- Objective function decomposes as a sum over training examples
 - This is unlike pure optimization
- Optimization for learning:
 - update the parameters based the expected value of cost using only a subset of the terms of of the full cost function

Ex: Decomposition into a sum

- Maximum likelihood estimation problem
 - In log-space estimated parameter decomposes into a sum over each example

$$\theta_{\text{ML}} = \arg \max_{\theta} \sum_{i=1}^m \log p_{\text{model}}(\mathbf{x}^{(i)}, y^{(i)}; \theta)$$

- It is equivalent to maximizing the expectation over the empirical distribution defined by the training set

$$J(\theta) = E_{\mathbf{x}, y \sim \hat{p}_{\text{data}}} \log p_{\text{model}}(\mathbf{x}, y; \theta)$$

- Commonly used property of $J(\theta)$ is its gradient

$$\tilde{N}_{\theta} J(\theta) = E_{\mathbf{x}, y \sim \hat{p}_{\text{data}}} \log p_{\text{model}}(\mathbf{x}, y; \theta)$$

- Computing this expectation is very expensive
 - Requires summation over every training sample
- Instead randomly sample small no. of samples

Quality of sampling-based estimate

- Standard error for mean from n samples is $\frac{\sigma}{\sqrt{n}}$
 - where σ is std dev of samples
- Denominator shows that error decreases less than linearly with no. of samples
 - Ex: 100 samples vs 10,000 samples
 - Computation increases by a factor of 100 but
 - Error decreases by only a factor of 10
- Optimization algorithms converge much faster
 - if allowed to rapidly compute approximate estimates
 - rather than slowly compute exact gradient

A motivation for sampling: Redundancy

- Training set may be redundant
 - Worst case: all m examples are identical
 - Sampling based estimate could use m times less computation
 - In practice
 - unlikely to find worst case situation but
 - likely to find large no. of examples that all make similar contribution to gradient

Batch gradient methods

- *Batch or deterministic gradient methods:*
 - Optimization methods that use all training samples in a large batch
- Somewhat confusing terminology
 - Batch also used to describe *minibatch* used by minibatch stochastic gradient descent
 - Batch gradient descent implies use of full training set
 - Batch size refers the size of a minibatch

Stochastic or online methods

- Those using a single sample are called *stochastic* or *on-line*
 - On-line typically means continually created samples, rather than multiple passes over a fixed size training set
- Deep learning algorithms use more than 1 but fewer than all
- Traditionally called *minibatch* or *minibatch stochastic* or simply *stochastic*

Minibatch Size

- Driven by following factors
 - Larger batches → more accurate gradient but with less than linear returns
 - Multicore architectures are underutilized by extremely small batches
 - Use some minimum size below which there is no reduction in time to process a minibatch
 - If all examples processed in parallel, amount of memory scales with batch size
 - This is a limiting factor in batch size
 - GPU architectures more efficient with power of 2
 - Range from 32 to 256, sometimes with 16 for large models

Regularizing effect of small batches

- Small batches offer regularizing effect due to noise added in process
- Generalization is best for batch size of 1
- Small batch sizes require small learning rate
 - To maintain stability due to high variance in estimate of gradient
- Total run time can be high
 - Due to reduced learning rate and
 - Requires more time to observe entire training set

Use of minibatch information

- Different algorithms use different information from the minibatch
 - Some algorithms more sensitive to sampling error
- Algorithms using gradient g are robust and can handle smaller batch sizes like 100
- Second order methods using Hessian H and compute updates such as $H^{-1} g$ require much larger batch sizes like 10,000

Random selection of minibatches

- Crucial to select minibatches randomly
- Computing expected gradient from a set of samples requires that sample independence
- Many data sets are arranged with successive samples highly correlated
 - E.g., blood sample data set has five samples for each patient
- Necessary to shuffle the samples
 - For a data set with billions of samples shuffle once and store it in shuffled fashion

Simple random sampling

- Define the population. Say, training set has 10,000 examples
- Choose your batch size: say 100
- List the population and assign numbers to them
- Use a random number generator to generate a number in $[1, 1000]$
- Select your sample

Parallelization of minibatches

- We can compute entire separate updates over different examples in parallel
 - Compute update that minimizes $J(X)$ for one minibatch of examples X at the same time we compute update for several other minibatches
- Synchronous parallel distributed approaches discussed in Section 12.1.3

SGD and generalization error

- Minibatch SGD follows the gradient of the true generalization error

$$J^*(\theta) = E_{(\mathbf{x}, y) \sim p_{data}} \left(L(f(\mathbf{x}; \theta), y) \right)$$

- As long as the examples are repeated
- Implementations of minibatch SGD
 - Shuffle once and pass through multiple no. of times
 - On the first pass: each minibatch computes unbiased estimate of true generalization error
 - Second pass: estimate is more biased because it is formed by resampling values already used rather than fair samples from data generating distribution

SGD minimizes generalization error

- Easiest to see equivalence in online learning
 - Examples/minibatches are drawn from a stream
 - Similar to living being
 - New example at each instant with each example (\mathbf{x}, y) coming from data generating distribution $p_{\text{data}}(\mathbf{x}, y)$
 - » Examples never repeated, every example is a fair sample
- Equivalence is easy to derive when \mathbf{x} and y are discrete
 - As seen next

Discrete case with loss function

- Generalization error (in terms of loss function)

is $J^*(\theta) = E_{(\mathbf{x}, y) \sim p_{data}} (L(f(\mathbf{x}; \theta), y))$

- Which can be written as a sum

$$J^*(\theta) = \sum_x \sum_y p_{data}(x, y) L(f(\mathbf{x}; \theta), y)$$

- with exact gradient

$$g = \nabla J^*(\theta) = \sum_x \sum_y p_{data}(x, y) \nabla_{\theta} L(f(\mathbf{x}; \theta), y)$$

Implies that derivative can be computed in batches

- We have already seen this (decomposition) demonstrated for log-likelihood in

$$J(\theta) = E_{\mathbf{x}, y \sim \hat{p}_{data}} \log p_{model}(\mathbf{x}, y; \theta) \quad \text{and} \quad \nabla_{\theta} J(\theta) = E_{\mathbf{x}, y \sim \hat{p}_{data}} \nabla_{\theta} \log p_{model}(\mathbf{x}, y; \theta)$$

- Thus it holds for functions other than likelihood
- Similar result for when \mathbf{x} and y are continuous

Use of multiple epochs

- SGD minimizes generalization error when samples are not reused
 - Yet best to make several passes through the training set
 - Unless training set is extremely large
- With multiple epochs, first epoch follows unbiased gradient of generalization error
- Additional epochs provide enough benefit to decrease training error
 - Although increasing gap between training and testing error

Impact of growing data sets

- Data sets are growing more rapidly than computing power
- More common to use each training example only once
 - Or even make an incomplete pass through the data set
- With a large training set overfit is not an issue
 - Underfitting and computational efficiency become predominant concerns