

Pooling in Convolutional Networks

Sargur Srihari
srihari@buffalo.edu

Topics in Convolutional Networks

- Overview
- 1. The Convolution Operation
- 2. Motivation
- 3. Pooling
- 4. Convolution and Pooling as an Infinitely Strong Prior
- 5. Variants of the Basic Convolution Function
- 6. Structured Outputs
- 7. Data Types
- 8. Efficient Convolution Algorithms
- 9. Random or Unsupervised Features
- 10. The Neuroscientific Basis for Convolutional Networks
- 11. Convolutional Networks and the History of Deep Learning

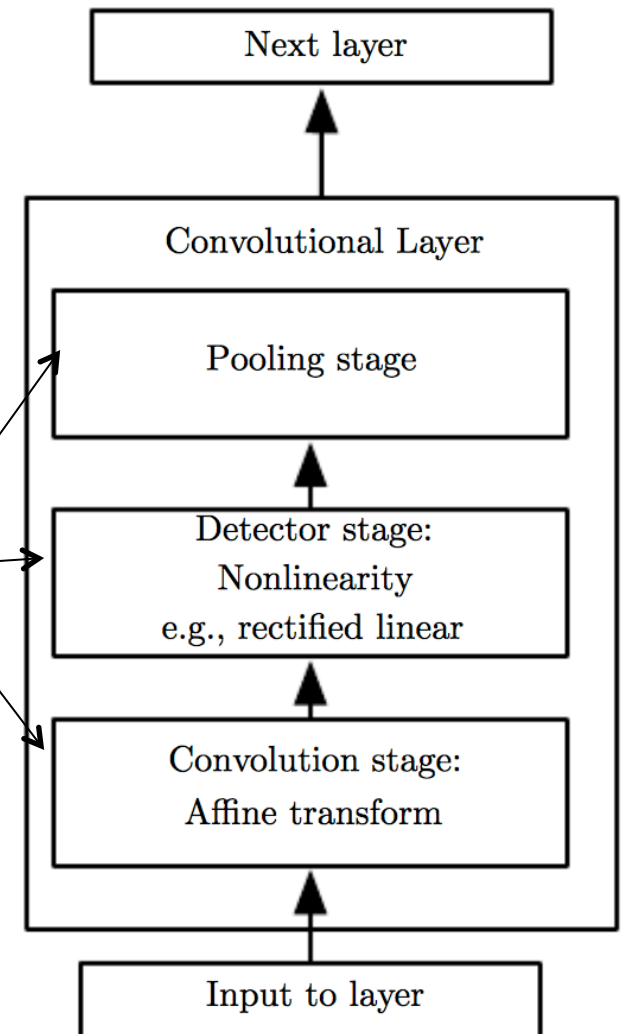
Topics in Pooling

- Three stages of CNNs
- Two terminologies: simple layers, complex layers
- Types of Pooling functions: Max, Average
- Translation invariance
- Rotation invariance
- Pooling with downsampling
- ConvNet Architectures

Pooling

- Typical layer of a CNN consists of three stages
- Stage 1:
 - perform several convolutions in parallel to produce a set of linear activations
- Stage 2 (Detector):
 - each linear activation is run through a nonlinear activation function such as ReLU
- Stage 3 (Pooling):
 - Use a pooling function to modify output of the layer further

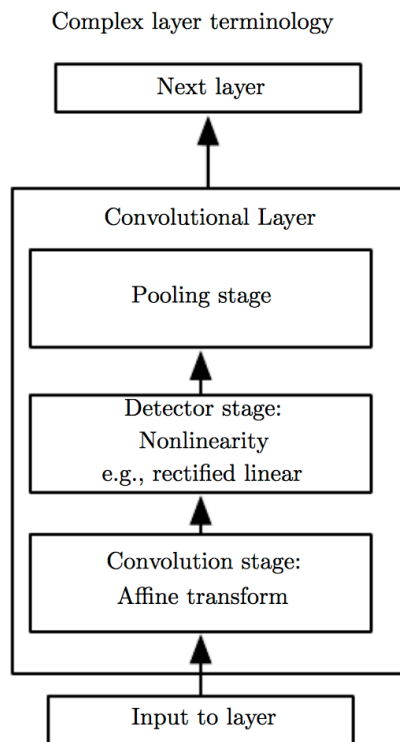
Complex layer terminology



Two terminologies for a typical CNN layer


1. Net is viewed as a small no. of complex layers, each layer having many stages

- 1-1 mapping between kernel tensors and network layers



2. Net is viewed as a larger no of simple layers

- Every processing step is a layer in its own right
- Not every layer has parameters



Types of Pooling functions

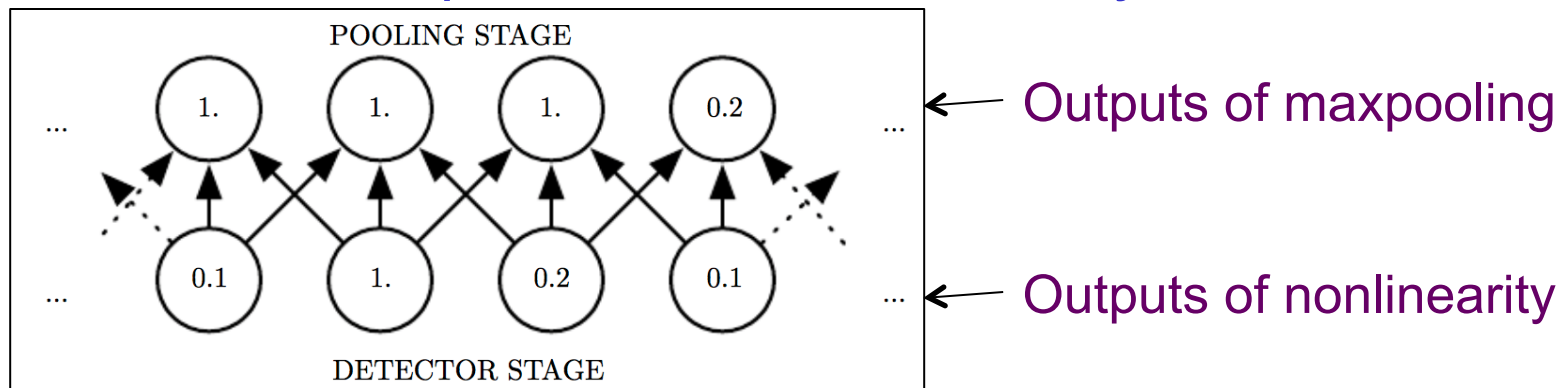
- A pooling function replaces the output of the net at a certain location with a summary statistic of the nearby inputs
- Popular pooling functions:
 1. *max pooling* operation reports the maximum output within a rectangular neighborhood
 2. Average of a rectangular neighborhood
 3. L^2 norm of a rectangular neighborhood
 4. Weighted average based on the distance from the central pixel

Pooling causes translation invariance


- In all cases, pooling helps make the representation become approximately *invariant* to small translations of the input
 - If we translate the input by a small amount values of most of the outputs does not change
 - Pooling can be viewed as adding a strong prior that the function the layer learns must be invariant to small translations

Max pooling introduces invariance to translation

- View of middle of output of a convolutional layer



- Same network after the input has been shifted by one pixel



- Every input value has changed, but only half the values of output have changed because maxpooling units are only sensitive to maximum value in neighborhood not exact value

Importance of Translation Invariance

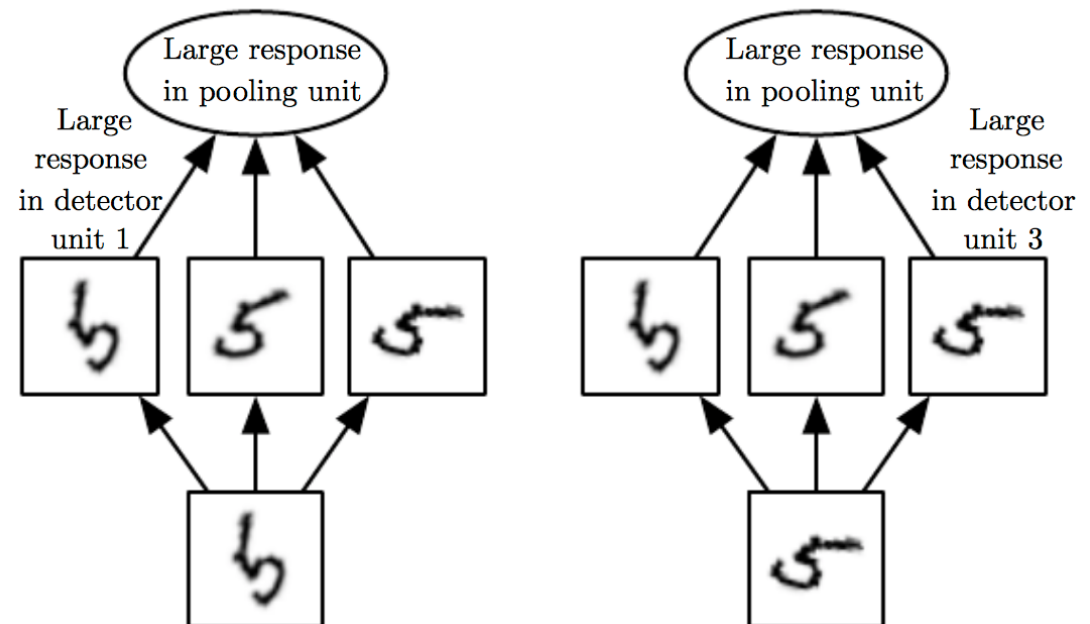
- Invariance to translation is important if we care about whether a feature is present rather than exactly where it is
 - For detecting a face we just need to know that an eye is present in a region, not its exact location
- In other contexts it is more important to preserve location of a feature
 - E.g., to determine a corner we need to know whether two edges are present and test whether they meet

Learning other invariances

- Pooling over spatial regions produces invariance to translation
- But if we pool over the results of separately parameterized convolutions, the features can learn which transformations to become invariant to

Learning Invariance to rotation

- A pooling unit that pools over multiple features that are learned with separate parameters can learn to be invariant to transformations of the input




Using fewer pooling units than detector units

- Because pooling summarizes the responses over a whole neighborhood, it is possible to use fewer pooling units than detector units
 - By reporting summary statistics for pooling regions spaced k pixels apart rather than one pixel apart
 - This improves computational efficiency because the next layer has k times fewer inputs to process
- An example is given next

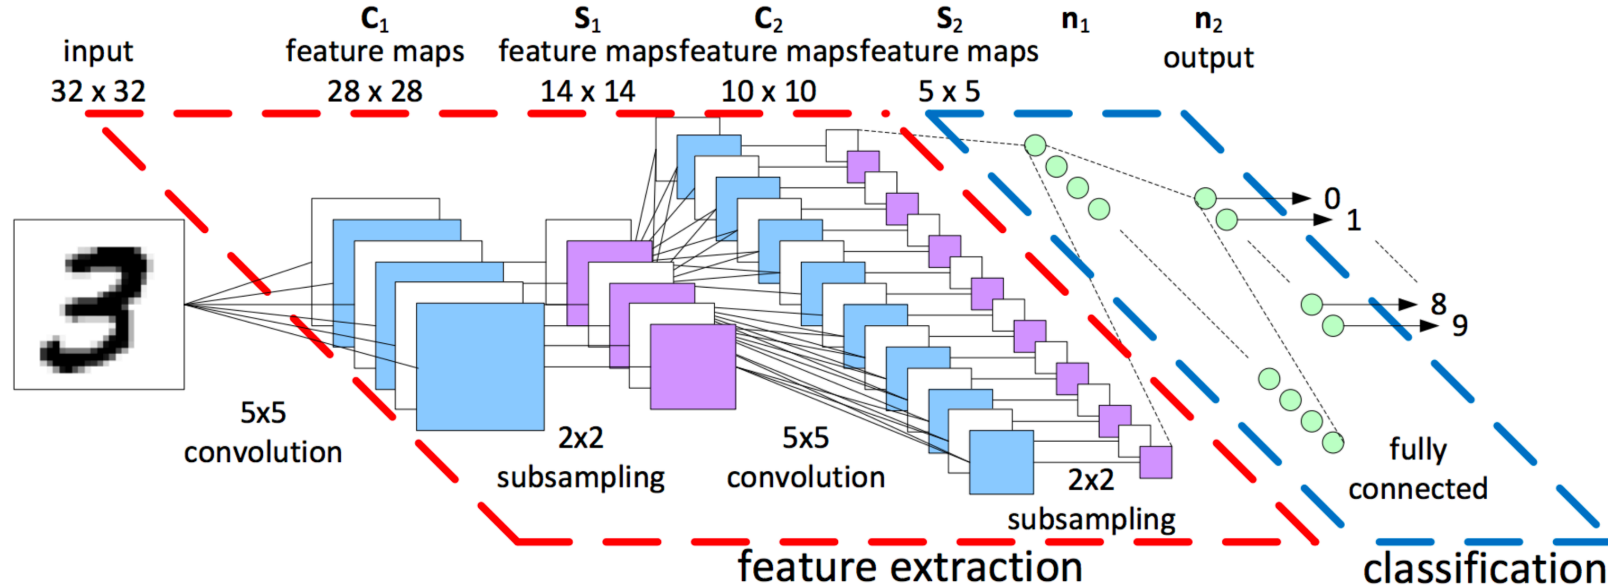
Pooling with down-sampling

- Max-pooling with a pool width of three and a stride between pools of two



- This reduces representation size by a factor of two
 - Which reduces computational burden of next layer
 - Rightmost pooling region has a smaller size but must be included if we don't want to ignore some of the detector units

Subsampling as Average pooling



Theoretical Guidance on Pooling

- Which kind of pooling one should use in different situations
- It is possible to dynamically pool features together
 - By running a clustering algorithm on locations of interesting features
 - Yields a different set of pooling regions for each image
 - Another approach: learn a single pooling structure
- Pooling can complicate architectures that use top-down information
 - E.g., Boltzmann machines and autoencoders

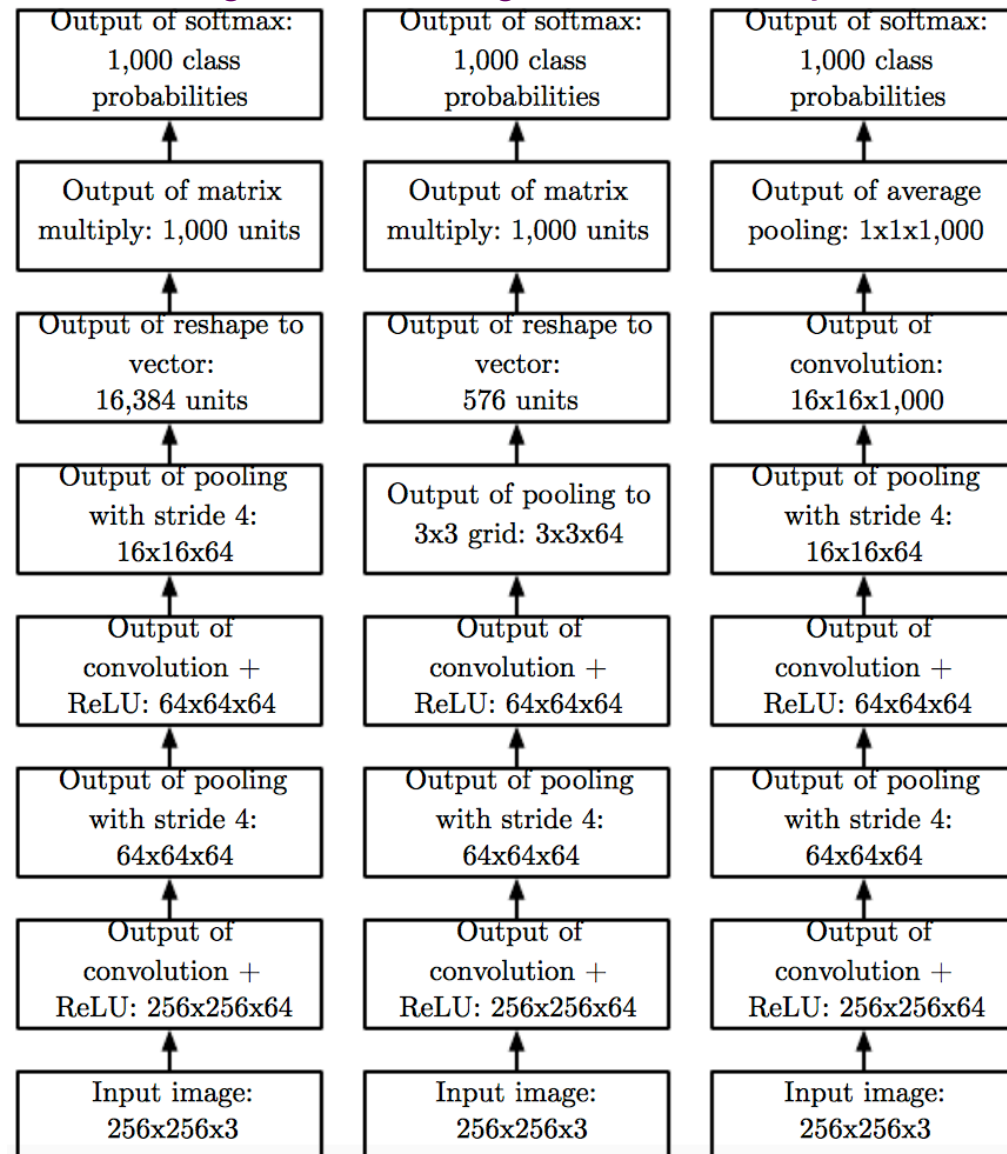
Examples of Architectures for Classification with CNNs

- Real networks have branching structures.
- Chain structures shown for simplicity

CNN that processes a fixed size image

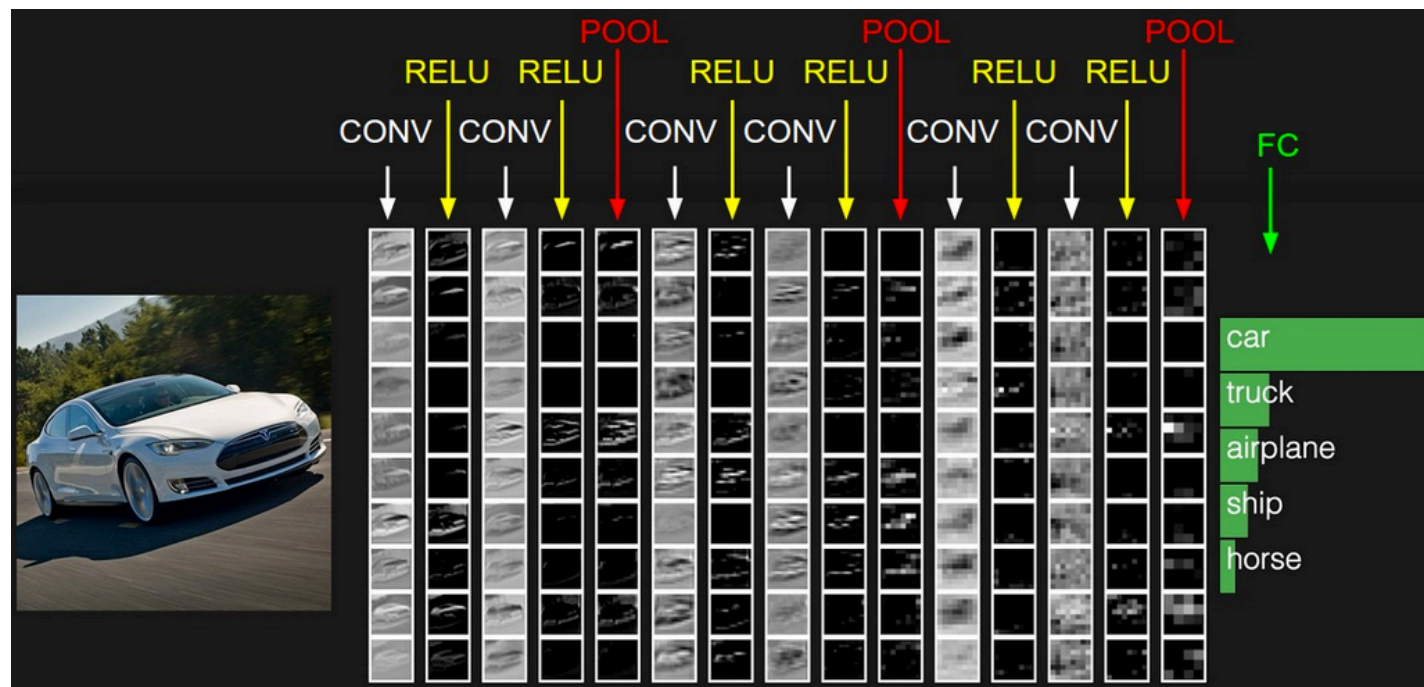
Processes variable-sized image

Does not have any fully-connected layer



A ConvNet architecture

INPUT: 32x32x3 holds raw pixel values: an image of width 32, height 32 and 3 color channels RGB
 CONV layer will compute the output of neurons connected to local regions in the input
 Each computing a dot product between their weights and a small region they are connected to
 In the input volume. This may result in a volume such as 32x32x12 if we used 12 filters
 POOL layer will perform a down-sampling operation along spatial dimensions (width, height)
 resulting in a volume such as 12x16x12



Activations of an example ConvNet architecture.

Initial volume stores raw image pixels (left) and the last volume stores class scores (right)


Each volume of activations along the processing path is shown as a column.

Since it is difficult to visualize 3D volumes, each volume's slices are laid out in rows

VGG Net

- VGG is a convolutional neural network model
 - K. Simonyan and A. Zisserman, University of Oxford
 - *“Very Deep Convolutional Networks for Large-Scale Image Recognition”*
- The model achieves 93% top-5 test accuracy in ImageNet
 - which is a dataset of over 14 million images belonging to 1000 classes.

VGG 16



Source: <https://www.cs.toronto.edu/~frossard/post/vgg16/>

VGG-16 pre-trained model for Keras

 vgg-16_keras.py

```

1 from keras.models import Sequential
2 from keras.layers.core import Flatten, Dense, Dropout
3 from keras.layers.convolutional import Convolution2D, MaxPooling2D, ZeroPadding2D
4 from keras.optimizers import SGD
5 import cv2, numpy as np
6
7 def VGG_16(weights_path=None):
8     model = Sequential()
9     model.add(ZeroPadding2D((1,1),input_shape=(3,224,224)))
10    model.add(Convolution2D(64, 3, 3, activation='relu'))
11    model.add(ZeroPadding2D((1,1)))
12    model.add(Convolution2D(64, 3, 3, activation='relu'))
13    model.add(MaxPooling2D((2,2), strides=(2,2)))
14
15    model.add(ZeroPadding2D((1,1)))
16    model.add(Convolution2D(128, 3, 3, activation='relu'))
17    model.add(ZeroPadding2D((1,1)))
18    model.add(Convolution2D(128, 3, 3, activation='relu'))
19    model.add(MaxPooling2D((2,2), strides=(2,2)))
20
21    model.add(ZeroPadding2D((1,1)))
22    model.add(Convolution2D(256, 3, 3, activation='relu'))
23    model.add(ZeroPadding2D((1,1)))
24    model.add(Convolution2D(256, 3, 3, activation='relu'))
25    model.add(ZeroPadding2D((1,1)))
26    model.add(Convolution2D(256, 3, 3, activation='relu'))
27    model.add(MaxPooling2D((2,2), strides=(2,2)))
28
29    model.add(ZeroPadding2D((1,1)))
30    model.add(Convolution2D(512, 3, 3, activation='relu'))
31    model.add(ZeroPadding2D((1,1)))

```

```

33    model.add(ZeroPadding2D((1,1)))
34    model.add(Convolution2D(512, 3, 3, activation='relu'))
35    model.add(MaxPooling2D((2,2), strides=(2,2)))
36
37    model.add(ZeroPadding2D((1,1)))
38    model.add(Convolution2D(512, 3, 3, activation='relu'))
39    model.add(ZeroPadding2D((1,1)))
40    model.add(Convolution2D(512, 3, 3, activation='relu'))
41    model.add(ZeroPadding2D((1,1)))
42    model.add(Convolution2D(512, 3, 3, activation='relu'))
43    model.add(MaxPooling2D((2,2), strides=(2,2)))
44
45    model.add(Flatten())
46    model.add(Dense(4096, activation='relu'))
47    model.add(Dropout(0.5))
48    model.add(Dense(4096, activation='relu'))
49    model.add(Dropout(0.5))
50    model.add(Dense(1000, activation='softmax'))
51
52    if weights_path:
53        model.load_weights(weights_path)
54
55    return model
56
57 if __name__ == "__main__":
58     im = cv2.resize(cv2.imread('cat.jpg'), (224, 224)).astype(np.float32)
59     im[:, :, 0] -= 103.939
60     im[:, :, 1] -= 116.779
61     im[:, :, 2] -= 123.68
62     im = im.transpose((2,0,1))
63     im = np.expand_dims(im, axis=0)
64
65     # Test pretrained model
66     model = VGG_16('vgg16_weights.h5')
67     sgd = SGD(lr=0.1, decay=1e-6, momentum=0.9, nesterov=True)
68     model.compile(optimizer=sgd, loss='categorical_crossentropy')
69     out = model.predict(im)
70     print np.argmax(out)

```

<https://gist.github.com/baraldilorenzo/07d7802847aaad0a35d3>