

# Recommender Systems

Sargur N. Srihari

srihari@cedar.buffalo.edu

This is part of lecture slides on [Deep Learning](http://www.cedar.buffalo.edu/~srihari/CSE676):  
<http://www.cedar.buffalo.edu/~srihari/CSE676>

# Topics in Recommender Systems

- Types of Recommender Systems
- Collaborative Filtering
- Word Embeddings to Item Embeddings
- Bilinear Prediction
- Relationship to Reinforcement Learning
  - Contextual Bandits
  - Exploration vs. Exploitation

Acknowledgements: Goodfellow, Bengio, Courville, Deep Learning, MIT Press, 2016

# Types of Recommender Apps

- A major family of applications of ML in the IT sector is the ability to make recommendations of items to potential users or customers
- Two major types of applications:
  - Online advertising
  - Item recommendations (for selling a product)
- Both rely on predicting association between user and an item

# Predicting user-item association

- Useful for either to predicting probability of some action
  - User buying product or some proxy for this action
- Or the expected gain (which may depend on the value of the product) if an ad is shown or a recommendation is made regarding that product to the user

# Commercial importance

- The internet is currently financed by various forms of online advertising
- Major parts of the economy rely on online shopping
- Amazon, eBay use ML including deep learning for product recommendations
- Sometimes items are not products for sale
  - E.g., selecting posts to display on social network feeds, recommending movies to watch, recommending jokes, recommending advise

# Collaborative Filtering

- Early work on recommender systems
  - Relied on minimal inputs for prediction
  - Rely on similarity between patterns of values of target variable for different users or different items
    - user 1 and user 2 both like items A, B and C
    - We may infer that user 1 and user 2 have similar tastes
    - If user 1 likes item D then this a strong cue that user 2 will also like D
    - Algorithms based on this principle come under the name of *collaborative filtering*

# Collaborative Filtering Methods

- Both non-parametric and parametric
- Non-parametric methods
  - Nearest neighbor based on similarity between patterns of preferences
- Parametric methods
  - Rely on learning a distributed representation called an *embedding* for each user and each item
  - Bilinear prediction of the target variable described next

# Word Embeddings to Item Embeddings

## Sentence-word data

s0: "Man walked his dog"

s1: "Man took his dog to park"

s2: "Dog went to park"

Vocabulary:

{man, walked, his, dog, took, to, park, went}

Vector representation:

s0: [1,1,1,1,0,0,0,0]


s1: [1,0,1,1,1,1,1,0]

s2: [0,0,0,1,0,1,1,1]

*Word2Vec* produces word embeddings in a low-dimensional continuous space and carry semantic and syntactic information of words

Continuous vector representations

Word2Vec



## User-item data

user0: Loc0, Loc1, Loc2, Loc3

user1: Loc0, Loc4, Loc2, Loc3, Loc5, Loc6

user2: Loc3, Loc7, Loc5, Loc6

Vocabulary:

{Loc0, Loc1, Loc2, Loc3, Loc4, Loc5, Loc6, Loc7}


Vector representation:

user0: [1,1,1,1,0,0,0,0]

user1: [1,0,1,1,1,1,1,0]

user2: [0,0,0,1,0,1,1,1]

Item2Vec



Similarly we get user2Vec



# Prediction of target variable

- Prediction of target variable (such as a rating)
  - Bilinear prediction is a simple parametric method
    - Highly successful
    - Found as a component in state-of-the-art systems
- Prediction is obtained by a dot product between the user embedding and the item embedding
  - Possibly corrected by by constants that depend only on either the user ID or item ID

# Bilinear Prediction

- Definitions

- Let  $R$  be the matrix containing our predictions
- $A$  a matrix with user embeddings in its rows
- $B$  a matrix with item embeddings in its columns
- Let  $b$  and  $c$  be vectors that contain respectively
  - a kind of bias for each user
    - Representing how grumpy or positive that user is
  - For each item
    - Representing its general popularity

- The bilinear prediction is  $\hat{R}_{u,i} = b_u + c_i + \sum_j A_{u,j} B_{j,i}$

- Typical goal: minimize squared error between predicted ratings  $\hat{R}_{u,i}$  and actual ratings  $R_{u,i}$

# Use of embeddings

- User embeddings and item embeddings can then be conveniently visualized when they are first reduced to a low dimension (two or three)
- Or they can be used to compare users or items against each other, just like word embeddings

# Obtaining embeddings

- One way to obtain these embeddings is by performing singular value decomposition (SVD) of the matrix  $R$  of actual targets (such as ratings)
- This corresponds to factorizing  $R=UDV'$  (or a normalized variant) into the product of two factors, the lower rank matrices  $A=UD$  and  $B=V'$

# Missing Entry problem with SVD

- One problem with SVD is that it treats missing entries in an arbitrary way, as if they corresponded to a target value of 0
- Instead we would like to avoid paying any cost for the predictions made on missing entries
- Fortunately the sum-of-squared-errors on the observed ratings can also be easily minimized by gradient-based optimization

# Netflix Competition

- Both SVD and bilinear prediction performed very well
- Competition was to predict ratings for films, based on previous ratings by a large set of anonymous users
- Even though it did not win by itself, the simple bilinear prediction of SVD was a component of the ensemble models presented by most competitors including the winners

$$\hat{R}_{u,i} = b_u + c_i + \sum_j A_{u,j} B_{j,i}$$

# Limitation of Collaborative Filtering

- When a new item or a new user is introduced, its lack of rating history means that there is no way to evaluate its similarity with other items or users (respectively),
- Or the degree of association between that user and existing items
- This is the problem of *cold-start recommendations*

# Solution to cold-start recommendation

- Introduce extra information about individual users or items
- Extra information could be user profile information or features of each item
- Systems that use such information are called *content-based recommender systems*
- The mapping from a rich set of user features or item features to an embedding can be learned through a deep learning architecture



# Content-based recommender systems

- Deep learning architectures such as CNNs learn to extract from rich content, e.g., musical audio tracks, for music recommendation
  - CNN takes acoustic features as input and computes an embedding for the associated song
  - Dot product between this song embedding and the embedding for a user is then used to predict whether a user will listen to the song

# Relationship to Reinforcement Learning

- A recommendation issue goes beyond supervised learning and into reinforcement learning
- Most recommendation problems are accurately described theoretically as *contextual bandits*
  - When recommendation system collects data, we get a biased and incomplete view of user preferences
    - We only see responses of users to items they were recommended and not to other items
    - In some cases no information on users for whom no recommendation has been made
      - E.g., with ad auctions, price proposed was below minimum, or does not win auction, so that ad is not shown

# Need for additional information

- System gets no information on what outcome would result from recommending any other item
  - It would be like training a classifier by picking one class  $\hat{y}$  for each training example  $x$  (typically class with highest probability) and then getting feedback whether this was correct or not
    - Each example conveys less information than in the supervised case where the true label  $y$  is directly observable, so more examples are necessary
      - We may keep on picking the wrong model output to show
    - The correct decision may have a low probability
      - Until learner picks the correct decision it does not learn about the correct decision

# Reinforcement Learning and Bandits

- In reinforcement learning only the reward for the selected action is observed
- In general, reinforcement learning can involve a sequence of many actions and many rewards
- Bandits scenario is a special case of reinforcement learning, in which the learner takes only a single action and receives a single reward
  - Bandit problem is easier in that the learner knows which reward is associated with which action

# Multi-armed Bandit Problem

1. Which machines to play
  2. How many times to play each machine
  3. In which order to play
- Each machine has
    - a probability distribution,  $B = \{R_1, \dots, R_K\}$ 
      - Mean values  $\mu_1, \dots, \mu_K$  associated with rewards
  - Objective:
    - Maximize reward through sequence of lever pulls
      - Minimize regret  $\rho$ , expected difference between optimal strategy and collected rewards  $r_t$ , after  $T$  rounds

One-armed  
bandit



Multi-armed  
bandit



$$\rho = T\mu^* - \sum_{t=1}^T \hat{r}_t$$

# Contextual Bandits

- In general reinforcement learning, a high or low reward might have been caused by a recent action or by an action in the distant past
- Contextual bandits refers to where the action is taken in context of an input variable that can inform the decision
  - E.g., we at least want to know user identity, and we want to pick an item
- Mapping from context to action is called policy
  - Feedback loop between learner and data distribution (which depends on actions of learner)
    - a central research issue in reinforcement learning

# Exploration vs Exploitation

- Reinforcement learning requires a tradeoff between exploration and exploitation
- Exploitation comes from taking actions that come from the current best version of the learned policy
  - Actions that we know will achieve a high reward
- Exploitation refers to taking actions specifically in order to obtain more training data

# Exploration example

- We know: in context  $x$ , action  $a$  has  $\text{reward}=1$ 
  - We don't know whether it's the best possible reward
  - We may want to exploit our current policy and continue taking action  $a$  in order to be relatively sure of obtaining a reward of 1
- However we may also want to explore action  $a'$ 
  - We do not know what will happen if we try action  $a'$
  - We hope for  $\text{reward}=2$  but we may get  $\text{reward}=0$
  - Either way we gain some knowledge



# Implementing Exploration-Exploitation

- Exploration can be implemented in many ways
  1. Random actions to cover all possible actions
  2. Model-based approaches
    - compute a choice of action based on its expected reward and the model's amount of uncertainty about that reward
- Factors determining exploration or exploitation
  - One prominent factor: Time scale
    - Agent has short time to accrue reward: exploitation
    - Agent has more time: begin with exploration so that future actions can be planned more effectively
      - As time progresses we move towards more exploitation

# Supervised Learning has no tradeoff

- No tradeoff between exploration-exploitation
- Supervision signal always specifies which output is correct for each input
- There is no need to try out different outputs to determine if one is better than the model's current output
  - We always know that the label is the best output

# Evaluating policies

- Another difficulty arising in the context of reinforcement learning
  - Besides exploration-exploitation tradeoff
- Difficulty of evaluating and comparing different policies
  - Reinforcement learning involves interaction between learner and environment
    - It is not straightforward to evaluate the learner's performance using a fixed set of test input values
    - The policy itself determines which inputs will be seen