

The Partition Function

Sargur N. Srihari
srihari@cedar.buffalo.edu

Topics

- Definition of Partition Function
 1. The log-likelihood gradient
 2. Stochastic maximum likelihood and contrastive divergence
 3. Pseudolikelihood
 4. Score matching and Ratio matching
 5. Denoising score matching
 6. Noise-contrastive estimation
 7. Estimating the partition function

Definition of Partition Function

- Undirected PGMs are defined by an unnormalized probability distribution $\tilde{p}(\mathbf{x}, \boldsymbol{\theta})$
 - which we must normalize by dividing by a *partition function* $Z(\boldsymbol{\theta})$ to obtain a valid probability distribution

$$p(\mathbf{x}, \boldsymbol{\theta}) = \frac{1}{Z(\boldsymbol{\theta})} \tilde{p}(\mathbf{x}, \boldsymbol{\theta})$$

- Where the partition function is either an integral or sum over unnormalized probabilities

$$Z(\boldsymbol{\theta}) = \int \tilde{p}(\mathbf{x}, \boldsymbol{\theta}) d\mathbf{x}$$

$$Z(\boldsymbol{\theta}) = \sum_x \tilde{p}(\mathbf{x}, \boldsymbol{\theta})$$

- This operation is intractable for many models

Dealing with Partition Functions

- Three approaches:
 1. Use a model with a tractable partition function
 2. Model is used in ways where $p(\mathbf{x})$ is not computed at all
 3. Directly confront the challenge of intractable partition functions
- We consider techniques for training and evaluating models with intractable partition functions

1. Log-likelihood gradient

- Learning undirected models is difficult because:
 - Partition function depends on parameters
- Gradient of the log-likelihood wrt parameters has a term corresponding to gradient of partition function

$$\nabla_{\theta} \log p(\mathbf{x}; \theta) = \nabla_{\theta} \log \tilde{p}(\mathbf{x}; \theta) - \nabla_{\theta} \log Z(\theta)$$

- Corresponds to positive and negative phase of learning

Tractability: Positive, Negative phases

- For most undirected models: negative phase is difficult
- Models with no latent variables or few interactions between latent variables have a tractable positive phase
- RBM: straight-forward positive phase, difficult negative phase
- Here we look at difficulties with the negative phase

Gradient of $\log Z$

$$\begin{aligned}\nabla_{\boldsymbol{\theta}} \log Z &= \frac{\nabla_{\boldsymbol{\theta}} Z}{Z} \\ &= \frac{\nabla_{\boldsymbol{\theta}} \sum_{\mathbf{x}} \tilde{p}(\mathbf{x})}{Z} \\ &= \frac{\sum_{\mathbf{x}} \nabla_{\boldsymbol{\theta}} \tilde{p}(\mathbf{x})}{Z}.\end{aligned}$$

- For models that guarantee $p(\mathbf{x}) > 0$ for all \mathbf{x} we can substitute $\exp(p_{\sim}(\mathbf{x}))$ for $p_{\sim}(\mathbf{x})$

$$\begin{aligned}& \frac{\sum_{\mathbf{x}} \nabla_{\boldsymbol{\theta}} \exp(\log \tilde{p}(\mathbf{x}))}{Z} \\ &= \frac{\sum_{\mathbf{x}} \exp(\log \tilde{p}(\mathbf{x})) \nabla_{\boldsymbol{\theta}} \log \tilde{p}(\mathbf{x})}{Z} \\ &= \frac{\sum_{\mathbf{x}} \tilde{p}(\mathbf{x}) \nabla_{\boldsymbol{\theta}} \log \tilde{p}(\mathbf{x})}{Z} \\ &= \sum_{\mathbf{x}} p(\mathbf{x}) \nabla_{\boldsymbol{\theta}} \log \tilde{p}(\mathbf{x}) = \mathbb{E}_{\mathbf{x} \sim p(\mathbf{x})} \nabla_{\boldsymbol{\theta}} \log \tilde{p}(\mathbf{x}).\end{aligned}$$

2. Stochastic Maximum Likelihood

- Naïve way of implementing

$$= \mathbb{E}_{\mathbf{x} \sim p(\mathbf{x})} \nabla_{\boldsymbol{\theta}} \log \tilde{p}(\mathbf{x}).$$

- Is to compute it by burning a set of Markov chains from a random initialization every time the gradient is needed
- When learning is performed by SGD this means chains must be burned in once per gradient step
- This approach leads to following training procedure

Naiive MCMC algorithm

- Maximizing log-likelihood with an intractable partition function using gradient ascent
-

Set ϵ , the step size, to a small positive number.

Set k , the number of Gibbs steps, high enough to allow burn in. Perhaps 100 to train an RBM on a small image patch.

while not converged **do**

Sample a minibatch of m examples $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$ from the training set.

$\mathbf{g} \leftarrow \frac{1}{m} \sum_{i=1}^m \nabla_{\boldsymbol{\theta}} \log \tilde{p}(\mathbf{x}^{(i)}; \boldsymbol{\theta})$.

Initialize a set of m samples $\{\tilde{\mathbf{x}}^{(1)}, \dots, \tilde{\mathbf{x}}^{(m)}\}$ to random values (e.g., from a uniform or normal distribution, or possibly a distribution with marginals matched to the model's marginals).

for $i = 1$ to k **do**

for $j = 1$ to m **do**

$\tilde{\mathbf{x}}^{(j)} \leftarrow \text{gibbs_update}(\tilde{\mathbf{x}}^{(j)})$.

end for

end for

$\mathbf{g} \leftarrow \mathbf{g} - \frac{1}{m} \sum_{i=1}^m \nabla_{\boldsymbol{\theta}} \log \tilde{p}(\tilde{\mathbf{x}}^{(i)}; \boldsymbol{\theta})$.

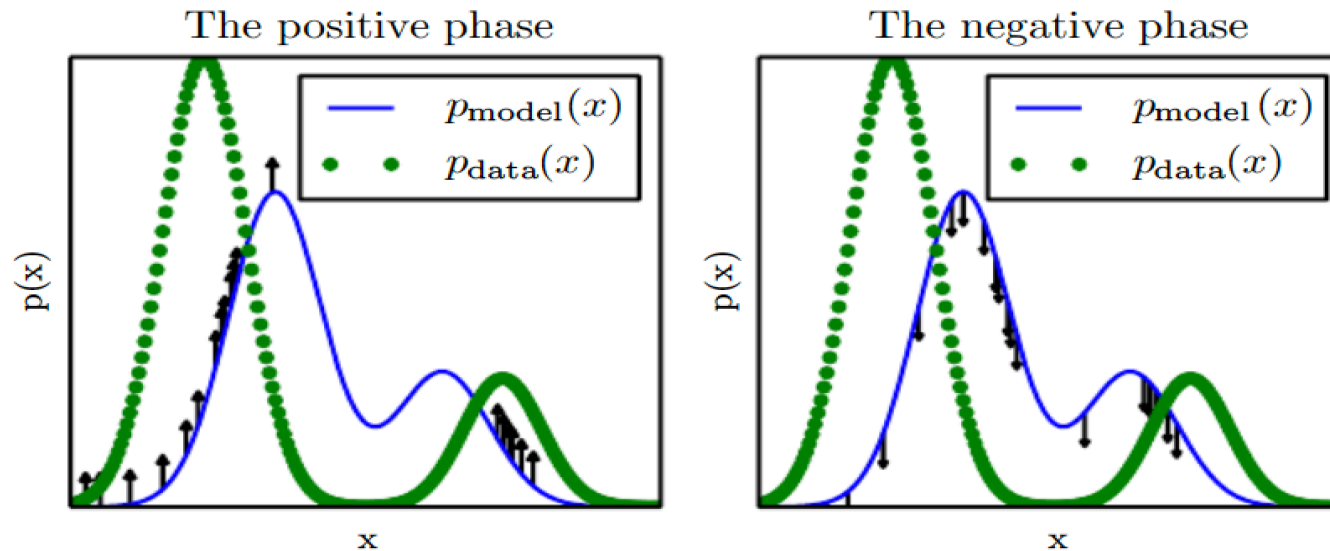
$\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} + \epsilon \mathbf{g}$.

end while

MCMC: balance between forces

- Can view MCMC approach as trying to balance between two forces
 - One pushing up on the model distribution where the data occurs
 - Another pushing down on the model distribution where the model samples occur
- Next figure illustrates this process
- The two forces correspond to
 - Maximizing $\log p_{\sim}$
 - Minimizing $\log Z$

Phases of MCMC Algorithm



Positive Phase:

We sample points from the data distribution and push up on their unnormalized probability

Negative Phase:

We sample points from the model distribution and push down on their unnormalized probability.

This counteracts positive phase's tendency to just add a constant to the unnormalized probability everywhere

When data distribution and model distribution are equal, the positive phase has the Same chance to push up at a point as the negative phase has to push down, When this happens, there is no longer any gradient (in expectation) and training must terminate.

Interpreting the negative phase

- Negative phase involves drawing samples from the model's distribution
 - Can think of it as finding points that the model believes in strongly
- Because the negative phase acts to reduce the probability of those points, they are generally considered to represent the model's incorrect beliefs about the world
- Referred to as *hallucinations* or *fantasy particles*

Neuroscientific analogy

- Negative phase proposed as explanation of dreaming in humans and other animals
 - Brain maintains a probabilistic model of the world
 - Follows the gradient of $\log p_{\sim}$ while experiencing real events while awake
 - Follows the negative gradient of $\log p_{\sim}$ to minimize $\log Z$ while sleeping and experiencing events sampled from the present model
- Not proven with neuroscientific experiments
 - In ML it is necessary to use positive and negative phases simultaneously
 - Rather than separate periods of wakefulness and REM sleep

A design less expensive than MCMC

- Given positive/negative phases of learning
 - We can design a less expensive alternative to naive MCMC
- Main cost of naive MCMC:
 - Cost of burning-in the Markov chains from a random initialization at each step
- Natural solution:
 - Initialize Markov chains from a distribution that is very close to the model distribution
 - So that burn-in operation does not take many steps

Contrastive Divergence algorithm

- Initializes the Markov chain at each step with samples from the data distribution
 - This is presented in algorithm given next
 - Obtaining samples from data distribution is free
 - Because they are already in the data set
- Initially data distribution is not close to model distribution, so the negative phase is inaccurate
 - Positive phase can increase model's data probability
 - After positive phase has time to act, the model distribution is closer to the data distribution
 - And the negative phase starts to become accurate

Contrastive Divergence Algorithm

- Using gradient ascent for optimization

Set ϵ , the step size, to a small positive number.

Set k , the number of Gibbs steps, high enough to allow a Markov chain sampling from $p(\mathbf{x}; \boldsymbol{\theta})$ to mix when initialized from p_{data} . Perhaps 1-20 to train an RBM on a small image patch.

while not converged **do**

 Sample a minibatch of m examples $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$ from the training set.

$\mathbf{g} \leftarrow \frac{1}{m} \sum_{i=1}^m \nabla_{\boldsymbol{\theta}} \log \tilde{p}(\mathbf{x}^{(i)}; \boldsymbol{\theta})$.

for $i = 1$ to m **do**

$\tilde{\mathbf{x}}^{(i)} \leftarrow \mathbf{x}^{(i)}$.

end for

for $i = 1$ to k **do**

for $j = 1$ to m **do**

$\tilde{\mathbf{x}}^{(j)} \leftarrow \text{gibbs_update}(\tilde{\mathbf{x}}^{(j)})$.

end for

end for

$\mathbf{g} \leftarrow \mathbf{g} - \frac{1}{m} \sum_{i=1}^m \nabla_{\boldsymbol{\theta}} \log \tilde{p}(\tilde{\mathbf{x}}^{(i)}; \boldsymbol{\theta})$.

$\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} + \epsilon \mathbf{g}$.

end while
