

# Learning as Optimization

Sargur Srihari

[srihari@cedar.buffalo.edu](mailto:srihari@cedar.buffalo.edu)

# Topics in Learning as Optimization

- Evaluation of Learned Model
- Empirical Risk and Overfitting
  - Bias vs. Variance Trade-off
  - Design and Evaluation of Learning Procedures
  - Goodness of Fit
  - PAC bounds
- Discriminative versus Generative Training
- Learning Tasks
  - Model Constraints
  - Data Observability
  - Taxonomy of Learning Tasks

# Optimization Viewpoint

- We have:
  1. Hypothesis Space
    - Set of candidate models

e.g., set of all BNs or MNs given a set of variables
  2. Objective Function
    - Criterion for quantifying preference over models
- Learning Task
  - Find a high-scoring model within model class
- Learning as optimization is predominant approach

# Criteria for Optimization

- Numerical Criteria (Loss functions) we optimize

## 1. Density Estimation

- Relative Entropy or K-L Divergence

- Expected value of log-difference

- Equivalent to Empirical Risk over instances  $\mathcal{D}$

- » Called Empirical log-loss

$$D(P^* \parallel P) = E_{\xi \sim P^*} \left[ \log \left( \frac{P^*(\xi)}{P(\xi)} \right) \right]$$

$$-\frac{1}{|\mathcal{D}|} \sum_{m=1}^M \log P(\xi[m] : \mathcal{M})$$

## 2. Classification

- Classification error: 0/1 loss

- Hamming loss

- Conditional log-likelihood

$$E_{(x,y) \sim P} [I\{h_p(x) \neq y\}]$$

$$E_{(x,y) \sim P^*} [\log P(y | x)]$$

- Can view learning as optimization

# Why discuss Optimization?

- Different choices of Objective Functions
  - Have ramification to results of learning procedures
- Has implications to further discussions on learning

# Empirical Distribution & Log-loss

- We have an unknown distribution  $P^*$

- Empirical Distribution

$$\hat{P}_D(A) = \frac{1}{M} \sum_m I\{\xi[m] \in A\}$$

Probability of Event  $A$  is the fraction of samples that satisfy  $A$

where  $\xi[1], \xi[2] \dots$  is a sequence of iid samples from  $P^*$

- For a sufficiently large training set  $P_D$ , will be quite close to  $P^*$

$$\lim_{M \rightarrow \infty} \hat{P}_{D_M}(A) = P^*(A)$$

- Consider empirical log-loss  $-\frac{1}{|\mathcal{D}|} \sum_{m=1}^M \log P(\xi[m] : \mathcal{M})$

- Can be shown that distribution that maximizes likelihood of data set  $D$  is the empirical distribution<sup>6</sup>

# Data Requirements

- How many data samples are needed?
- Consider two cases
  1. 100 binary random variables
  2. Bayesian network where a node has  $k$  parents
- In both we will see that the hypothesis space is too large

# Data Requirement with Binary Variables

- Consider  $100$  binary variables
  - $2^{100}$  possible joint assignments
- $\mathcal{D}$  has  $1,000$  instances (most likely distinct)
  - probability of  $0.001$  to each assignment
  - $0$  to remaining  $2^{100}-1000$
- Example is extreme, but phenomenon is general



# Data Requirement with Bayesian network

- $M^*$  is a Bayesian network with a variable 'Fever'
  - “Fever” has many parents (diseases)  $X_1, \dots, X_k$
  - In a table CPD
    - parameters grow exponentially with no. of parents  $k$
  - Unlikely  $\mathcal{D}$  has cases representing all parent instantiations
    - i.e., all combinations of diseases  $X_1, \dots, X_k$
    - Results in very poor CPDs
- Data requirement: exponential with BN connectivity
  - No of samples needed grows linearly with no. of parameters
  - No. of parameters grows exponentially with network connectivity

# Problem with using Empirical Risk

- Learned model tends to overfit to training data
- But our goal is to answer queries about unseen data
  - Not data samples we have seen
    - New patients in disease example
    - New unsegmented images
- Need to generalize well
- This leads to a trade-off
  - Should not overfit but generalize well

# Overfitting and Bias-Variance Trade-off

- If hypothesis space is very limited (few bins), unable to represent true  $P^*$ 
  - Even with unlimited data
  - Called Bias
- If hypothesis space is large (expensive)
  - Examples: binary data set and disease BN
  - If data set is small even random fluctuations in data set have an effect
  - Results in large variance over different data sets

# Implication to Model Structure

- Not allow too rich a class of models
- With limited data
  - Error due to variance  $\geq$  Error due to bias
  - Restrict to models too simple to correctly encode  $P^*$
  - Ability to estimate parameters reliably
- compensates for
  - Error from incorrect structural assumptions
- Restricting models  $\rightarrow$  learn important edges
  - Distance to  $P^*$  is better
- Combining approaches
  - Hard constraint over model class
  - Optimization objective that leads us away from over-fitting

# Generalization Performance

- Have not discussed
  - How well model generalizes
    - Performance on unseen data sets
  - Design of hypothesis space to reduce over-fitting
- Next:
  1. Basic experimental protocol
    - in design/evaluation of learning procedure
  2. Theoretical framework for appropriate complexity

# Relative and Absolute Learning

## 1. Relative

- Compare two or more alternatives
  - Different hypothesis spaces
  - Different training objectives
- Methods: *Holdout* and *Cross-Validation*
  - Training, Validation, Testing

## 2. Absolute

- Whether model captures true distribution
- Method: *Goodness of Fit*

# Evaluating Generalization Performance

- Performance on unseen data
- Hold-out method
  - Partition data into two sets:  $\mathcal{D}_{train}$ ,  $\mathcal{D}_{test}$
  - Use randomized procedure to decide partition
  - Learn models using  $\mathcal{D}_{train}$  and an objective function
  - Measure performance using  $\mathcal{D}_{test}$  and loss function
  - Provides unbiased estimate on new samples
- Performance on  $\mathcal{D}_{train}$  usually better than  $\mathcal{D}_{test}$ 
  - If difference is large then we are overfitting
    - Use less expressive model
    - or method to discourage overfitting

# $k$ -fold Cross Validation

- If data set is small
- Divide into  $k$  sets
- Test on one while training on the remaining
- For writing a paper on learning
  - Report results using cross-validation
- In practice
  - Train on all data



# Selecting a Learning Procedure

- Validation set
- If we are testing among a several procedures
  - Using the same testing set to select the procedure yields an over-optimistic result
    - Since selection is optimized to  $\mathcal{D}_{test}$
- Use three sets
  - $\mathcal{D}_{train}$ : used to learn the model
  - $\mathcal{D}_{validate}$ : evaluate different variants of training
  - $\mathcal{D}_{test}$ : evaluate final performance

# Goodness of Fit

- Whether learned model captures everything about the distribution
  - Harder to answer
- After determining the parameters we have a hypothesis about distribution
  - Now ask whether data behave as if they were generated from this distribution

# Implementation of Goodness of Fit

- Let  $f$  be a property of a data set
  - E.g., empirical log-loss
- Evaluate  $f(\mathcal{D}_{train})$
- From learned  $\mathcal{M}$  generate new data sets  $\mathcal{D}$
- If  $f(\mathcal{D}_{train})$  deviates significantly from distribution of  $f(\mathcal{D})$ 
  - Reject hypothesis that  $\mathcal{D}_{train}$  was generated from  $\mathcal{M}$

# Goodness of Fit: Choosing $f$

- Natural choice: Empirical log-loss in data set

Expected Loss

$$E_D[\text{loss}(\xi : M)] = \frac{1}{|\mathcal{D}|} \sum_{\xi \in \mathcal{D}} \text{loss}(\xi : \mathcal{M}) \qquad \frac{1}{|\mathcal{D}|} \sum_{m=1}^M \log P(\xi[m] : \mathcal{M})$$

- Test if log-loss for  $\mathcal{D}_{train}$  differs significantly from expected empirical log-loss (on  $\mathcal{D}$  sampled from  $\mathcal{M}$ )
  - Expected value is simply the entropy of  $\mathcal{M}$ 
    - Entropy of BN can be computed efficiently
- To check for significance
  - Consider tail distribution of log-loss— more involved
  - Approximation:
    - variance of log-loss as function of  $M$
  - Alternatively: generate large no of data sets  $\mathcal{D}$  of size  $M$  from model and estimate distribution over  $E_D[\text{loss}(\xi : \mathcal{M})]$

# Probably Approximately Correct

- Given target loss function we can estimate empirical risk on training set  $\mathcal{D}_{train}$
- Because of over-fitting to  $\mathcal{D}_{train}$ 
  - may not reflect performance on new data  $\mathcal{D}_{new}$
- Are the two performances ( $\mathcal{D}_{train}$  &  $\mathcal{D}_{new}$ ) related?
  - Does low training loss imply low expected loss?
- We cannot guarantee with certainty quality of learned model
  - Data set  $\mathcal{D}$  is sampled stochastically from  $P^*$
  - Sample may be unrepresentative
    - Probability is low but not zero

# PAC Bound

- Assume relative entropy loss function
- $P_M^*$ : distribution over data sets  $\mathcal{D}$  of size  $M$ 
  - sampled from  $P^*$
- Learning Procedure  $L$  given  $\mathcal{D}$  returns model  $\mathcal{M}_{L(\mathcal{D})}$
- We want to prove results of the form

Let  $\varepsilon$  be our approximation parameter

$\partial > 0$  our confidence parameter

For  $M$  large enough we have

$$P_M^* (\{ \mathcal{D} : D(P^* \| P_{\mathcal{M}_{L(\mathcal{D})}}) \leq \varepsilon \}) \geq 1 - \partial$$

No of samples  $M$   
needed to reach bound  
is called is called  
sample complexity

Type of result is  
called a PAC bound

- i.e., for sufficiently large  $M$  for most data sets of size  $M$  sampled from  $P^*$ , learning procedure  $L$  applied to  $\mathcal{D}$  will learn a close approximation to  $P^*$

# PAC Bound in Practice

- Can only be obtained if hypothesis space contains the model to correctly represent  $P^*$
- Many cases hypothesis space may not contain  $P^*$ 
  - Cannot expect to learn a model whose relative entropy to  $P^*$  is guaranteed to be low
- Best we can hope is to get a model whose error is at most  $\epsilon$  worse than the lowest error found in hypothesis space
  - Expected loss beyond minimum possible error is called excess risk


# Discriminative vs Generative Training

- We assumed that goal is to learn model  $\mathcal{M}$  to be a good approximation for  $P^*$
- Often we want the model to perform well on a given task
  - E.g., predicting  $Y$  from  $X$
- Generative training objective:
  - To get  $\mathcal{M}$  close to  $P^*(Y, X)$
- Discriminative Training objective:
  - To get  $P(Y|X)$  close to  $P^*(Y|X)$
- Same model class trained in these two ways can get different results



# Example of Generative/Discriminative Training

- Markov network with
  - Target variable  $Y$ , Features  $X_1, \dots, X_n$
  - $\Phi_i(X_i, Y)$ ,  $\Phi_0(Y)$
- If trained generatively to fit  $P(Y, X_1, \dots, X_n)$  we are learning a naïve Markov model



– Which is equivalent to Naïve Bayes

- Since network is singly connected

$$P(Y, X_1, \dots, X_k) = P(Y) \prod_{i=1}^k P(X_i | Y)$$

- Pairwise clique potentials look like CPDs

- If trained discriminatively, to fit  $P(Y | X_1, \dots, X_n)$ , we are learning the logistic regression model

$$P(Y = 1 | x_1, \dots, x_k) = \text{sigmoid} \left\{ w_0 + \sum_{i=1}^k w_i x_i \right\}$$

# Discrim. Training Disadvantages

- Generative model still useful for prediction
  - E.g., Naïve Bayes used for classification
- Discriminative model trained for prediction  $P(Y|X)$  does not encode distribution over  $X$ 
  - No conclusion about these variables
- Discriminative training less appealing for BNs
  - DT of naïve Bayes gives pairwise potentials locally normalized CPDs
- Discriminative training mostly used for CRFs


# Generative Models have higher Bias

- Generative models make more assumptions about form of distribution
  - Encode independences about feature variables  $X$
  - Discriminative models make independence assumptions only about  $Y$  and their dependency on  $X$
- Generative model defines  $P(Y, X)$ 
  - Thereby induces  $P(Y|X)$  and  $P(X)$  using same overall model
    - Thus training needs to fit both well
  - Discriminative model needs a good fit only for  $P(Y|X)$
- With limited sets of data generative model less likely to over-fit
  - As data set grows discriminative model less likely to be affected by model assumptions

# Example: OCR problem

$Y$  = character label ( $A, B, \dots$ )

$X_i$  are pixels with 8-bit values ( $1, \dots, 256$ )



- Generatively train naïve Markov
  - Naïve Markov (or Bayes) learns distribution over 256 values given each label  $A, B, \dots$
  - Each estimated independently
    - Low-dimensional estimation problem
- Discriminatively train logistic regression
  - Jointly optimizes  $26 \times 256$  parameters
    - Much higher dimensional estimation problem
- For sparse data naïve Bayes is better
- Discriminative Classifier can use much richer features
  - Edges, image patch centered at pixel which use same pixels

{ Pixel independence  
Is false assumption

$$P(Y = 1 | x_1, \dots, x_k) = \text{sigmoid} \left\{ w_0 + \sum_{i=1}^k w_i x_i \right\}$$

{ Pixel correlation  
accounted for

# Learning Tasks

- Different Variants of the learning task
- Vary along three axes:
- Two types of Input to learning procedure
  1. Prior knowledge about constraints about  $\mathcal{M}$ 
    - Know graph structure, learn parameters
    - Learn both structure and parameters
    - Not know full set of variables over which  $P^*$  is defined
  2. Data  $\mathcal{D} = \{d[1], \dots, d[M]\}$  which are i.i.d. samples of  $P^*$ 
    - Fully observed
    - Partially observed
      - E.g., Patient records: not all tests performed, Disease not fully certain
    - Hidden variables
      - Their inclusion can simplify structure, e.g., genetic susceptibility to a disease
- Output is model  $\mathcal{M}$ 
  3. May include structure and parameters
- Define a Taxonomy of Learning Tasks

# Parameter Estimation for Known Structure

## Numerical Optimization

- Numbers are difficult to elicit from people
- Parameter Estimation is basis for more advanced scenarios

- Bayesian networks  
with complete data

- Parameter estimation is easily solved
- Possible closed-form solution

- Markov networks

- Global partition function induces entanglements of parameters
- For fixed structure problems
  - Optimization problem is convex
    - Iterative numerical optimization
    - Each step requires inference: expensive

Convex minimization :

- (i) local minimum is global minimum.
- (ii) set of all (global) minima is convex.
- (iii) For *strictly* convex minimum is unique.

# Approaches to Structure Learning

- Constraint-based Structure Learning
- Score-based Structure Learning
- Bayesian Model Averaging

# Constraint-based Structure Learning

- View Bayesian network as a representation of independencies
- Sensitive to failures of individual independence tests
- If one test returns a wrong answer it misleads the network construction procedure



# Score-based Structure Learning

- View a BN as specifying a statistical model
- Learning is a model selection problem
  - Hypothesis space is superexponential ( $2^{O(n^2)}$ )
    - Situation is worse for Markov nets since cliques can be of size greater than two
  - Each network structure is given a score
    - Optimize to find highest score
- Search problem may not have an elegant and efficient solution

# Bayesian Model Averaging

- Generates an ensemble of possible structures
- Average the prediction of all possible structures
- Due to immense number of structures, approximations are needed

# Dealing with Incomplete Data

- Multiple hypotheses regarding values of unobserved variables lead to
- Combinatorial range of alternative models
- Induce a non-convex, multimodal optimization problem in parameter space
- Known algorithms work iteratively
  - Using EM style algorithms